

Anirudha Acharya P

+91 9964577004 | anirudhaacharyap@gmail.com | linkedin.com/in/anirudhaacharyap | github.com/anirudhaacharyap

EDUCATION

Nitte Meenakshi Institute of Technology

B.E., Artificial Intelligence & Machine Learning — GPA: 7.93/10

Bengaluru, India

2023 – 2027

EXPERIENCE

Google Student Ambassador

Google

Aug 2025 – Jan 2026

- Drove AI/ML adoption across NMIT, reaching 1,000+ students through hands-on developer workshops
- Co-organized two large-scale vibe coding events with Google and Vercel; managed end-to-end event logistics for 500+ attendees per event

PROJECTS

Vibe App | *Java, Spring Boot, Kafka, Redis, PostgreSQL, gRPC, Python, Flutter*

- Architected a distributed event-driven backend processing 5,000+ events/day over Kafka with zero message loss under sustained load
- Cut API p95 latency by 30% with Redis caching; exposed music-taste ML microservice over gRPC achieving sub-20ms inference round-trip

AlloyAdvisor | *Python, FastAPI, React, TypeScript, LLM APIs*

- Built a full-stack AI agent for metal alloy recommendations with dual chat/structured-query interface; Finished top 10 Bangalore Codex Hackathon out of 80+ teams
- Designed real-time conflict detection engine + multi-turn LLM tool-use pipeline, cutting query resolution from 3 lookup steps to one interaction

Multi-Source Scraping & Analytics Platform | *Python, FastAPI, PostgreSQL, Docker*

- Built fault-tolerant ingestion pipelines across YouTube, Google Trends, and forums with auto-retry on rate limits; composable API layer powering 50+ workflows

Colorectal Cancer Classification | *Python, PyTorch, Diffusion Transformers*

- Designed hybrid BOA-WOA optimizer + Diffusion Transformer achieving AUC-ROC 1.0000, MCC 0.9960, and 95% CI lower bound 0.9950 — exceeding published SOTA benchmarks on both metrics
- Proved hybrid superiority over single-component baselines, reducing misclassifications from 6 to 2 per 1,000 samples

OPEN SOURCE

Evidently — ML Observability & Data Drift Library

- Contributed to data drift detection and report generation modules used by production ML teams globally

XMem — The Memory Layer for AI That Never Forgets

- Implemented POST /v1/memory/batch-ingest endpoint for batch conversation ingestion, reducing network overhead by eliminating per-turn HTTP requests
- Added strict request/response schemas, reused existing ingestion pipeline, and wrote unit tests covering success workflows and schema validation

TECHNICAL SKILLS

- **Languages:** Python, Java, TypeScript, JavaScript, C++, Go
- **Backend:** FastAPI, Spring Boot, Node.js, REST APIs, gRPC, Kafka, GraphQL
- **Frontend:** React, TypeScript
- **Databases:** Oracle, PostgreSQL, Redis, SQL, MongoDB
- **Infrastructure:** AWS, Docker, Git, Linux, CI/CD
- **AI / LLM:** Agentic workflows, LLM tool-use, RAG pipelines, Prompt engineering, N8N
- **Concepts:** DSA, Computer Networks, OS, DBMS, Software Development, System Design, Distributed Systems

ACHIEVEMENTS

- Finished top 10 at Bangalore Codex Hackathon (AlloyAdvisor) among 80+ competing teams
- Founded two student clubs at NMIT: Oasis (Gaming & Game Dev) and Diseno Divino (UI/UX & AI)